

CS 331: Artificial Intelligence

Local Search II

1

3. Beam Search

2

Local Beam Search

Travelling Salesman Problem




Keeps track of k states rather than just 1. $k=2$ in this example. Start with k randomly generated states.

3

Local Beam Search Example

Travelling Salesman Problem ($k=2$)




Generate all successors of all the k states

4

Local Beam Search Example

Travelling Salesman Problem ($k=2$)




None of these is a goal state so we continue

5

Local Beam Search Example

Travelling Salesman Problem ($k=2$)



Select the best k successors from the **complete** list

Local Beam Search Example

Travelling Salesman Problem ($k=2$)



Repeat the process until goal found

7

Local Beam Search

- How is this different from k random restarts in parallel?
- Random-restart search: each search runs independently of the others
- Local beam search: useful information is passed among the k parallel search threads
- Eg. One state generates good successors while the other $k-1$ states all generate bad successors, then the more promising states are expanded

8

Local Beam Search

- Disadvantage: all k states can become stuck in a small region of the state space
- To fix this, use stochastic beam search
- Stochastic beam search:
 - Doesn't pick best k successors
 - Chooses k successors at random, with probability of choosing a given successor being an increasing function of its value

9

4. Genetic Algorithms

10

Genetic Algorithms




- Like natural selection in which an organism creates offspring according to its fitness for the environment
- Essentially a variant of stochastic beam search that combines two parent states (just like sexual reproduction)
- Over time, population contains individuals with high fitness

11

Definitions


- **Fitness function:** Evaluation function in GA terminology
- **Population:** k randomly generated states (called individuals)
- **Individual:** String over a finite alphabet



12

Definitions

- **Selection:** Pick two random individuals for reproduction
- **Crossover:** Mix the two parent strings at the crossover point



Definitions

- **Mutation:** randomly change a location in an individual's string with a small independent probability

4	8	1	5	1	6	2	3	4	2
---	---	---	---	---	---	---	---	---	---



4	8	1	5	1	6	0	3	4	2
---	---	---	---	---	---	---	---	---	---

Randomness aids in avoiding small local extrema

14

GA Overview

Population = Initial population
 Iterate until some individual is fit enough or enough time has elapsed:
 NewPopulation = Empty
 For 1 to size(Population)
 Select pair of parents (P_1, P_2) using Selection(P, Fitness Function)
 Child C = Crossover(P_1, P_2)
 With small random probability, Mutate(C)
 Add C to NewPopulation
 Population = NewPopulation
 Return individual in Population with best Fitness Function

15

GA Overview

Population = Initial population
 Iterate until some individual is fit enough or enough time has elapsed:
 NewPopulation = Empty
 For 1 to size(Population)
 Select pair of parents (P_1, P_2) using Selection(P, Fitness Function)
 Child C = Crossover(P_1, P_2)
 With small random probability, Mutate(C)
 Add C to NewPopulation
 Population = NewPopulation
 Return individual in Population with best Fitness Function

This pseudocode only produces one child. Could also do a variant like before where we produce 2 children

16

Lots of variants


- Variant 1: Culling - individuals below a certain threshold are removed
- Variant 2: Selection based on:

$$P(X \text{ selected}) = \frac{Eval(X)}{\sum_{Y \in Population} Eval(Y)}$$

17

Example: 8-queens

- Fitness Function: number of nonattacking pairs of queens (28 is the value for the solution)
- Represent 8-queens state as an 8 digit string in which each digit represents position of queen



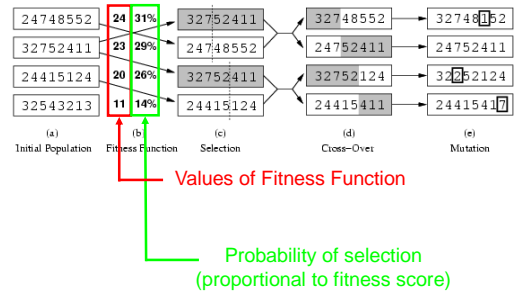
18

Example: 8-queens




19

Example: 8-queens (Fitness Function)



20



Example: 8-queens (Selection)



Notice 3 2 7 5 2 4 1 1 is selected twice while
3 2 5 4 3 2 1 3 is not selected at all

21

Example: 8-queens (Crossover)



22

Example: 8-queens (Mutation)



Mutation corresponds to randomly selecting a queen and randomly moving it in its column

23

Implementation details on Genetic Algorithms

- Initially, population is diverse and crossover produces big changes from parents
- Over time, individuals become quite similar and crossover doesn't produce such a big change
- Crossover is the big advantage:
 - Preserves a big block of "genes" that have evolved independently to perform useful functions
 - E.g. Putting first 3 queens in positions 2, 4, and 6 is a useful block

24

Schemas

- A substring in which some of the positions can be left unspecified eg. 246*****
- Instances: strings that match the schema
- If the average fitness of the instances of a schema is above the mean, then the number of instances of the schema within the population will grow over time

25

Schemas

- Schemas are important if contiguous blocks provide a consistent benefit
- Genetic algorithms work best when schemas correspond to meaningful components of a solution

26

The fine print...

- The representation of each state is critical to the performance of the GA
- Lots of parameters to tweak but if you get them right, GAs can work well
- Limited theoretical results (skeptics say it's just a big hack)

27

And remember....

```
def getSolutionCosts(navigationCode):
    fuelStopCost = 15
    extraComputationCost = 8
    thisAlgorithmBecomingSkynetCost = 999999999
    waterCrossingCost = 415
```

GENETIC ALGORITHMS TIP:
ALWAYS INCLUDE THIS IN YOUR FITNESS FUNCTION

(From <http://www.xkcd.com/534/>)

28

Gradient Descent

29

Discrete Environments

Hillclimbing pseudocode

```
X ← Initial configuration
Iterate:
    E ← Eval(X)
    N ← Neighbors(X)
    For each Xi in N
        Ei ← Eval(Xi)
    E* ← Highest Ei
    X* ← Xi with highest Ei
    If E* > E
        X ← X*
    Else
        Return X
```

- In discrete state spaces, the # of neighbors is finite.
- What if there is a continuum of possible moves leading to an infinite # of neighbors?

30

Local Search in Continuous State Spaces

- Almost all real world problems involve continuous state spaces
- To perform local search in continuous state spaces, you need techniques from calculus
- The main technique to find a minimum is called **gradient descent** (or **gradient ascent** if you want to find the maximum)

31

Gradient Descent

- What is the gradient of a function $f(x)$?
 - Usually written as

$$\nabla f(x) = \frac{\partial}{\partial x} f(x)$$
 - $\nabla f(x)$ (the gradient itself) represents the direction of the steepest slope
 - $|\nabla f(x)|$ (the magnitude of the gradient) tells you how big the steepest slope is

32

Gradient Descent

Suppose we want to find a local minimum of a function $f(x)$. We use the gradient descent rule:

$$x \leftarrow x - \alpha \nabla f(x)$$

α is the learning rate, which is usually a small number like 0.05


Suppose we want to find a local maximum of a function $f(x)$. We use the gradient ascent rule:

$$x \leftarrow x + \alpha \nabla f(x)$$



33

Gradient Descent Examples

$$f(x) = x^3 - 2x^2 + 2$$



These pictures were taken from Wolfram Mathworld



34

Question of the Day

- Why not just calculate the global optimum using $\nabla f(x) = 0$?
 - May not be able to solve this equation in closed form
 - If you can't solve it globally, you can still compute the gradient locally (like we are doing in gradient descent)

35

Multivariate Gradient Descent

- What happens if your function is multivariate eg. $f(x_1, x_2, x_3)$?
- Then

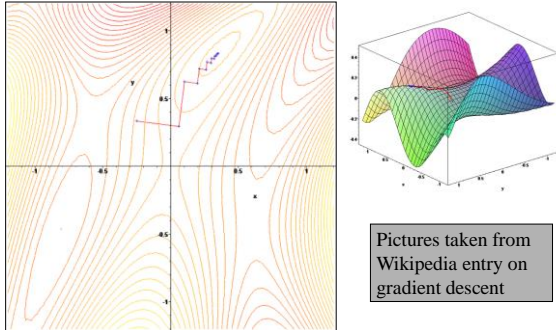
$$\nabla f(x_1, x_2, x_3) = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial x_3} \right)$$
- The gradient descent rule becomes:

$$x_1 \leftarrow x_1 - \alpha \frac{\partial f}{\partial x_1} \quad x_3 \leftarrow x_3 - \alpha \frac{\partial f}{\partial x_3}$$

$$x_2 \leftarrow x_2 - \alpha \frac{\partial f}{\partial x_2}$$

36

Multivariate Gradient Ascent



37

More About the Learning Rate

- If α is too large
 - Gradient descent overshoots the optimum point
- If α is too small
 - Gradient descent requires too many steps and will take a very long time to converge

38

Weaknesses of Gradient Descent

1. Can be very slow to converge to a local optimum, especially if the curvature in different directions is very different
2. Good results depend on the value of the learning rate α
3. What if the function $f(x)$ isn't differentiable at x ?

39

What you should know

- Be able to formulate a problem as a Genetic Algorithm
- Understand what crossover and mutation do and why they are important
- Differences between hillclimbing, simulated annealing, local beam search, and genetic algorithms
- Understand how gradient descent works, including its strengths and weaknesses
- Understand how to derive the gradient descent rule

40