




- Everyone needs to login to os-class and cleanup lingering message queues from the /dev/mqueue directory. There are many old message queues in the directory that are consuming kernel resources.
- If you have multiple posixmsg_server processes running, you are very likely to run into problems. Kill off your old processes. It is easy and makes your life better.
- **Exit functions are good things.**



Sockets: Internet Domain

Internet domain **stream** sockets are implemented on top of TCP. They provide a reliable, bidirectional, byte-stream communication channel.

Internet domain datagram sockets are implemented on top of UDP.



TCP: Transmission Control Protocol. TCP provides reliable, ordered, and error-checked delivery of a stream of octets between applications running on hosts communicating over an IP network.

UDP: User Datagram Protocol. UDP uses a simple connectionless transmission model with a minimum of protocol mechanism. It has no handshaking dialogues, and thus exposes any unreliability of the underlying network protocol to the user's program. There is no guarantee of delivery, ordering, or duplicate protection. UDP provides checksums for data integrity, and port numbers for addressing different functions at the source and destination of the datagram. Time-sensitive applications often use UDP because dropping packets is preferable to waiting for delayed packets, which may not be an option in a real-time system.

Network Byte Order

	2-byte integer		4-byte integer			
	address	address	address	address	address	address
	N	$N + 1$	N	$N + 1$	$N + 2$	$N + 3$
Big-endian byte order	1 (MSB)	0 (LSB)	3 (MSB)	2	1	0 (LSB)
Little-endian byte order	0 (LSB)	1 (MSB)	0 (LSB)	1	2	3 (MSB)

MSB = Most Significant Byte, LSB = Least Significant Byte

R. Jesse Chaney

CS344 – Oregon State University

One problem we encounter when passing these data across a network is that different hardware architectures store the bytes of a multi-byte integer in different orders.

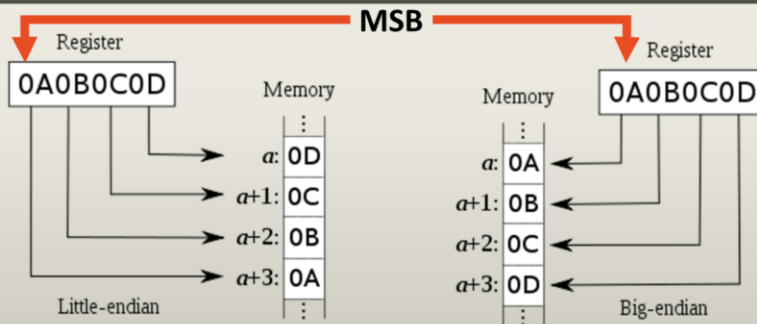
I'm sure you've talked about big-endian and little-endian in other classes.

The [Internet Protocol](#) defines big-endian as the standard *network byte order* used for all numeric values in the [packet headers](#) and by many higher level protocols and file formats that are designed for use over IP. The [Berkeley sockets API](#) defines a set of functions to convert 16-bit and 32-bit integers to and from network byte order: the `htons` (host-to-network-short) and `htonl` (host-to-network-long) functions convert 16-bit and 32-bit values respectively from machine (*host*) to network order; the `ntohs` and `ntohl` functions convert from network to host order. These functions may be a [no-op](#) on a big-endian system.

In 1726, [Jonathan Swift](#) described in his satirical novel [Gulliver's Travels](#) tensions in [Lilliput and Blefuscu](#): whereas royal edict in Lilliput requires cracking open one's [soft-boiled egg](#) at the small end, inhabitants of the rival kingdom of Blefuscu crack theirs at the big end (giving them the moniker *Big-endians*).^{[6][7]} The terms *little-endian* and

endianness have a similar intent.

Network Byte Order



R. Jesse Chaney

CS344 – Oregon State University

Big-endian is the most common convention in data networking (including IPv6), hence its pseudo-synonym **network byte order**, and little-endian is popular (though not universal) among microprocessors in part due to **Intel's** significant historical influence on microprocessor designs.

Failure to account for varying endianness across architectures when writing software code for mixed platforms and when exchanging certain types of data might lead to failures and bugs, though these issues have been understood and properly handled for many decades.



Ordering Your Bytes



```
#include <arpa/inet.h>

uint16_t htons (uint16_t host_uint16 );
    Returns host_uint16 converted to network byte order
uint32_t htonl (uint32_t host_uint32 );
    Returns host_uint32 converted to network byte order

uint16_t ntohs (uint16_t net_uint16 );
    Returns net_uint16 converted to host byte order
uint32_t ntohl (uint32_t net_uint32 );
    Returns net_uint32 converted to host byte order
```

R. Jesse Chaney

CS344 – Oregon State University

Notice that a short is 16 bits and a long is 32 bits.



The C long data type may be 32 bits on some systems and 64 bits on others.

With structures, the issue is further complicated by the fact that different implementations employ different rules for **aligning** the fields of a structure to address boundaries on the host system, leaving different numbers of padding bytes between the fields.

Just when you thought it would be easy to switch back and forth between host-order and network order, a whole new set of complications come up.

Because of these differences in data representation, applications that exchange data between heterogeneous systems over a network must adopt some common convention for encoding that data. The sender must encode data according to this convention, while the receiver decodes following the same convention. The process of putting data into a standard format for transmission across a network is referred to as marshalling.

A simpler approach than marshalling is often employed: encode all transmitted data in text form, with separate data items delimited by a designated character, typically a newline character. One advantage of this approach is that we can use telnet to debug an application.

In the book, the author goes into some detail about how to write a `newline()` function that will return a NULL terminated string from a socket.



Internet Socket Addresses

There are two types of Internet domain socket addresses: IPv4 and IPv6.

IPv6 address is 128 bits, much larger than the 32 bits of IPv4.

IPv4 address exhaustion



R. Jesse Chaney

CS344 – Oregon State University

If IPv4 and IPv6 coexist on a host, they share the same port-number space. This means that if, for example, an application binds an IPv6 socket to TCP port 2000 (using the IPv6 wildcard address), then an IPv4 TCP socket can't be bound to the same port.

In this class, we'll mostly deal with the IPv4 addresses. However, the world is moving to the IPv6 standard, so the more you know about that, the better off you are.

IPv4 address exhaustion is the depletion of the pool of unallocated Internet Protocol Version 4 (IPv4) addresses, which has been anticipated since the late 1980s. This depletion is the reason for the development and deployment of its successor protocol, IPv6.

IPv4 provides approximately 4.3 billion addresses.

On 31 January 2011, the last two unreserved IANA /8 address blocks were allocated to APNIC according to RIR request procedures. This left five reserved but unallocated /8 blocks.

IPv6 uses a 128-bit address, allowing 2^{128} , or approximately 3.4×10^{38} addresses, or more than 7.9×10^{28} times as many as IPv4. IPv6 addresses are represented as eight groups of four hexadecimal digits separated by colons, for example

2001:0db8:85a3:0042:1000:8a2e:0370:7334

, but methods to abbreviate this full notation exist.

How Big is IPv6?



- 4.295×10^9 **IPv4** address space size (32-bit)
- 10^{22} Estimated number of stars in the universe.
- 4.339×10^{26} Nanoseconds that have passed since the Big Bang.
- 3.156×10^{31} Unique IPv6 addresses that would be assigned after 1 trillion years if a new IPv6 address was assigned at every picosecond.
- 3.403×10^{38} **IPv6** address space size (128-bit).
- 10^{81} Estimated total number of atoms in the universe



Host and Service Conversion Functions



The **inet_aton()** and **inet_ntoa()** functions convert an **IPv4** address in dotted-decimal notation to binary and from binary to dotted-decimal.

Use these!

The **inet_pton()** and **inet_ntop()** functions are like **inet_aton()** and **inet_ntoa()**, but differ in that they **also handle IPv6 addresses**.

R. Jesse Chaney

CS344 – Oregon State University

Computers represent IP addresses and port numbers in binary. However, humans find names easier to remember than numbers. Employing symbolic names also provides a useful level of indirection; users and programs can continue to use the same name even if the underlying numeric value changes.

A hostname is the symbolic identifier for a system that is connected to a network (possibly with multiple IP addresses). A service name is the symbolic representation of a port number.

inet_pton() and **inet_ntop()** convert **binary IPv4 and IPv6** addresses to and from presentation format—that is, either dotted-decimal or hexstring notation.



```
#include <arpa/inet.h>
```

```
int inet_pton(int domain  
    , const char *src_str ,void *addrptr );
```

Returns 1 on successful conversion, 0 if src_str is not in presentation format,
or -1 on error

```
const char *inet_ntop(int domain  
    , const void * addrptr  
    , char *dst_str, size_t len );
```

Returns pointer to dst_str on success, or NULL on error

The p in the names of these functions stands for “presentation,” and the n stands for “network.”

Host and Service Conversion

```
#include <sys/socket.h>
#include <netdb.h>

int getaddrinfo(const char * host
    , const char * service
    , const struct addrinfo * hints
    , struct addrinfo ** result );
    Returns 0 on success, or nonzero on error
```



The **getaddrinfo()** function converts host and service names to IP addresses and port numbers. It was defined in POSIX.1g as the successor to the **obsolete** **gethostbyname()** and **getservbyname()** functions.

R. Jesse Chaney

CS344 – Oregon State University

The **getnameinfo()** function is the converse of **getaddrinfo()**. It translates a socket address structure (either IPv4 or IPv6) to strings containing the corresponding host and service name.

The **getaddrinfo()** is a bit more awkward to use than the older functions, but provides greater flexibility.

As output, **getaddrinfo()** dynamically allocates a linked list of **addrinfo** structures and sets **result** pointing to the beginning of this list.

The **getaddrinfo()** function dynamically allocates memory for all of the structures referred to by **result** (Figure 59-3). Consequently, the caller must deallocate these structures when they are no longer needed. The **freeaddrinfo()** function is provided to conveniently perform this deallocation in a single step.



ifconfig

\$ ifconfig

```
eth0 Link encap:Ethernet HWaddr 74:86:7A:E0:FD:B4
      inet addr:128.193.37.168 Bcast:128.193.39.255 Mask:255.255.252.0
      inet6 addr: fe80::7686:7aff:ee0:fdb4/64 Scope:Link
      UP BROADCAST RUNNING MULTICAST MTU:1500 Metric:1
      RX packets:831636402 errors:0 dropped:82 overruns:0 frame:0
      TX packets:705209123 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:1000
      RX bytes:380621753171 (354.4 GiB) TX bytes:660091655852 (614.7 GiB)
      Interrupt:35

lo    Link encap:Local Loopback
      inet addr:127.0.0.1 Mask:255.0.0.0
      inet6 addr: ::1/128 Scope:Host
      UP LOOPBACK RUNNING MTU:65536 Metric:1
      RX packets:469184 errors:0 dropped:0 overruns:0 frame:0
      TX packets:469184 errors:0 dropped:0 overruns:0 carrier:0
      collisions:0 txqueuelen:0
      RX bytes:561981920 (535.9 MiB) TX bytes:561981920 (535.9 MiB)
```

This is probably
what you are
looking for.

R. Jesse Chaney

CS344 – Oregon State University

You can also try `ip addr`, but I don't find it easier to use.