

# Bayes and Naïve Bayes Classifiers

CS434

# In this lecture

1. Review some basic probability concepts
2. Introduce a useful probabilistic rule - Bayes rule
3. Introduce the learning algorithm based on Bayes rule (thus the name Bayes classifier) and its extension, Naïve Bayes

# Commonly used discrete distributions

Binary random variable  $x \sim \text{Bernoulli}(p)$

$$\left. \begin{array}{l} P(x = 1) = p; \\ P(x = 0) = 1 - p \end{array} \right\} p(x) = p^x (1 - p)^{(1-x)}$$



Categorical distribution:  $x$  can take multiple values,  $v_1, \dots, v_k$

$$\left. \begin{array}{l} P(x = v_1) = p_1 \\ P(x = v_2) = p_2 \\ P(x = v_k) = p_k \end{array} \right\} p(x) = \prod_{i=1}^k p_i^{I(x=v_i)}$$



$$p_1 + p_2 + \dots + p_k = 1$$

# Learning the parameters of discrete distributions

- Let  $x$  denote the event of getting a head when tossing a coin ( $x \in \{0,1\}$ )
- Given a sequence of  $n$  coin tosses  $x_1, \dots, x_n$ , we estimate

$$p(x = 1) = \frac{1}{n} \sum_{i=1}^n x_i$$

- Let  $x$  denote the outcome of rolling a die ( $x \in \{1, \dots, 6\}$ )
- Given a sequence of  $n$  rolls, we estimate

$$\begin{aligned} p(x = j) \\ &= \frac{1}{n} \sum_{i=1}^n I(x_i = j) \end{aligned}$$


Estimation using such simple counting performs what we call “Maximum Likelihood Estimation” (MLE). There are other methods as well.

# Continuous Random Variables

- A continuous random variable  $x$  can take any value in an interval on the real line
  - $x$  usually corresponds to some real-valued measurements, e.g.,  $x = \text{today's lowest temperature}$
  - The probability of a continuous random variable taking an exact value is typically zero:  $P(x=56.2)=0$
  - It is more meaningful to measure the probability of a random variable taking a value within an interval  $P(x \in [50, 60])$
  - This is captured in the ***Probability density function***

# PDF: probability density function

- We often use  $f(x)$  to denote the PDF of  $x$
- $f(x) \geq 0$
- $f(x)$  can be larger than 1
- $\int_{-\infty}^{\infty} f(x)dx = 1$
- $\int_{x_1}^{x_2} f(x)dx = P(x_1 < x < x_2)$




- If  $f(x_1) = \alpha f(x_2)$ :

When  $x$  is sampled from  $f(x)$ , you are  $\alpha$  times as likely to see a  $x$  value “near”  $x_1$  than that a  $x$  value “near”  $x_2$

# Commonly Used Continuous Distributions

## Uniform Probability Density Function


$$f(x) = \begin{cases} 1/(b-a) & \text{for } a \leq x \leq b \\ 0 & \text{elsewhere} \end{cases}$$



E.g., Suppose we know that in the last 10 minutes, one customer arrived at OSU federal counter #1. The actual time of arrival X can be modeled by a uniform distribution over the interval of (0, 10)

## Normal (Gaussian) Probability Density Function

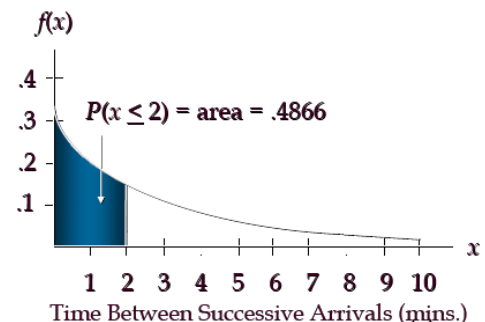
$$f(x) = \frac{1}{\sqrt{2\pi}\sigma} e^{-(x-\mu)^2/2\sigma^2}$$



E.g., the body temp. of a person, the average IQ of a class of 1<sup>st</sup> graders

## Exponential Probability Distribution

$$f(x) = \frac{1}{\mu} e^{-x/\mu}$$




E.g., the time between two successive forest fires

# Conditional Probability

- $P(A|B)$  = probability of A being true given that B is true

H = "Have a headache"

F = "Coming down with Flu"



$$P(H) = 1/10$$


$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

"Headaches are rare and flu is rarer, but if you're coming down with a flu there's a 50-50 chance you'll have a headache."



# Conditional Probability



H = "Have a headache"  
F = "Coming down with  
Flu"

$$P(H) = 1/10$$

$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

$$P(H|F)$$

$$= \frac{\text{Area of "H and F" region}}{\text{Area of "F" region}}$$

$$= \frac{P(H \wedge F)}{P(F)}$$


# Definition of Conditional Probability

$$P(A|B) = \frac{P(A \wedge B)}{P(B)}$$

## Corollary: The Chain Rule

$$P(A \wedge B) = P(A|B) P(B)$$

# Probabilistic Inference



H = "Have a headache"

F = "Coming down with Flu"

$$P(H) = 1/10$$


$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

One day you wake up with a headache. You think: "Drat! 50% of flus are associated with headaches so I must have a 50-50 chance of coming down with flu"

Is this reasoning good?

# Probabilistic Inference



H = "Have a headache"

F = "Coming down with Flu"

$$P(H) = 1/10$$

$$P(F) = 1/40$$

$$P(H|F) = 1/2$$

$$P(F \wedge H) = \dots$$

$$P(F|H) = \dots$$

# What we just did...


$$P(B|A) = \frac{P(A \wedge B)}{P(A)} = \frac{P(A|B) P(B)}{P(A)}$$

This is Bayes Rule


**Bayes, Thomas (1763)** An essay towards solving a problem in the doctrine of chances. *Philosophical Transactions of the Royal Society of London*, **53:370-418**



# Using Bayes Rule to Gamble




The “Win” envelope  
has a dollar and four  
beads in it




The “Lose” envelope  
has three beads and  
no money

Trivial question: someone draws an envelope at random and offers to sell it to you. How much should you pay in order to not lose money on average?

# Using Bayes Rule to Gamble



The “Win” envelope  
has a dollar and four  
beads in it



The “Lose” envelope  
has three beads and  
no money

Interesting question: before deciding, you are allowed to see one bead drawn from the envelope.

Suppose it's black: How much should you pay?

Suppose it's red: How much should you pay?

# Where are we?

- We have recalled the fundamentals of probability
- We have discussed conditional probability and Bayes rule
- Now we will move on to talk about the Bayes classifier



# Basic Idea

- Each example is described by  $m$  input features, i.e.,  $\mathbf{x} = [x_1, x_2, \dots, x_m]^T$ , for now we assume they are discrete variables
- Each example belongs to one of  $c$  possible classes, denoted by  $y \in \{1, \dots, c\}$
- If we don't know anything about the features, a randomly drawn example has a fixed probability  $P(y = j)$  to belong to class  $j$
- If an example belongs to class  $j$ , its features  $\mathbf{x} = [u_1, \dots, u_m]^T$  will follow some particular distribution  $p(u_1, \dots, u_m | y = j)$
- Given an example  $\mathbf{x} = [u_1, \dots, u_m]^T$ , a bayes classifier reasons about the value of  $y$  using Bayes rule:

$$P(y = j | u_1, \dots, u_m) = \frac{P(u_1, \dots, u_m | y = j)P(y = j)}{P(u_1, \dots, u_m)}$$

# Learning a Bayes Classifier

$$P(y = j | u_1, \dots, u_m) = \frac{P(u_1, \dots, u_m | y = j)P(y = j)}{P(u_1, \dots, u_m)}$$

Given a set of training data  $S = \{(\mathbf{x}_i, y_i) : i = 1, \dots, n\}$ , we learn:

- **Class Priors:**  $P(y = j)$  for  $j = 1, \dots, c$ 
$$P(y = j) = \frac{1}{n} \sum_{i=1}^n I(y_i = j)$$
- **Class Conditional Distribution:**  $P(\mathbf{x} = [u_1, \dots, u_m]^T | y = j)$  for  $j = 1, \dots, c$ 
$$P(u_1, \dots, u_m | y = j) = \frac{\text{\# of class } j \text{ examples that have values } (u_1, u_2, \dots, u_m)}{\text{\# of total class } j \text{ examples}}$$
- **Marginal Distribution of  $\mathbf{x}$ :**  $P(u_1, \dots, u_m)$

No need to learn, can be computed using

$$P(\mathbf{x}) = \sum_{j=1}^c P(\mathbf{x} | y = j)P(y = j)$$

# Learning a joint distribution

Build a JD table for your attributes in which the probabilities are unspecified

A	B	C	Prob
0	0	0	?
0	0	1	?
0	1	0	?
0	1	1	?
1	0	0	?
1	0	1	?
1	1	0	?
1	1	1	?

Fraction of all records in which  
A and B are True but C is False

The fill in each row with

$$\hat{P}(\text{row}) = \frac{\text{records matching row}}{\text{total number of records}}$$

A	B	C	Prob
0	0	0	0.30
0	0	1	0.05
0	1	0	0.10
0	1	1	0.05
1	0	0	0.05
1	0	1	0.10
1	1	0	<b>0.25</b>
1	1	1	0.10

# Example of Learning a Joint

- This Joint was obtained by learning from three attributes in the UCI “Adult” Census Database [Kohavi 1995]

48842 examples in total

12363 examples with this value combination

gender	hours_worked	wealth	
Female	v0:40.5-	poor	0.253122
		rich	0.0245895
	v1:40.5+	poor	0.0421768
		rich	0.0116293
Male	v0:40.5-	poor	0.331313
		rich	0.0971295
	v1:40.5+	poor	0.134106
		rich	0.105933

UCI machine learning repository:  
<http://archive.ics.uci.edu/ml/datasets/Adult>

# Bayes Classifiers in a nutshell

1. Estimate  $P(y = j)$  as fraction of class  $j$  examples for  $j=1,\dots,c$ .
  2. Learn conditional joint distribution  $p(x_1, x_2, \dots, x_m \mid y = j)$  for  $j = 1, \dots, c$
  3. To make a prediction for a new example  $\mathbf{x} = [u_1, \dots, u_m]^T$
- learning*

$$\begin{aligned} y^{\text{predict}} &= \underset{j}{\operatorname{argmax}} P(y = j \mid u_1, \dots, u_m) \\ &= \underset{j}{\operatorname{argmax}} \frac{P(u_1, \dots, u_m \mid y = j) P(y = j)}{P(u_1, \dots, u_m)} \\ &= \underset{j}{\operatorname{argmax}} P(u_1, \dots, u_m \mid y = j) P(y = j) \end{aligned}$$

If you wish to have the probability value  $P(y = 1 \mid [u_1, \dots, u_m]^T)$ , you will need to compute the normalizing factor  $P([u_1, \dots, u_m]^T)$ .

# Example: Spam Filtering

- Use the Bag-of-words representation
- Assume a dictionary of  $m$  words and tokens
- Create one binary attribute for each dictionary entry
  - i.e.,  $x_i = 1$  means the  $i$ -th word is used in the email
- Consider a reasonable dictionary size:  $m = 5000$  --- we have 5000 binary attributes
- How many parameters that we need to learn?
  - We need to learn the class prior  $P(y=1), P(y=0)$ : 1
  - For each of the two classes, we need to learn a joint distribution table of 5000 binary variables:  $2^{5000} - 1$
  - **Total:  $2 * (2^{5000} - 1) + 1$**
- Clearly we don't have nearly enough data to estimate that many parameters

# Joint Distribution Overfits

- It is common to encounter the following situation:
  - No training examples have the exact value combinations  $\mathbf{x} = (u_1, u_2, \dots, u_m)$ ,
  - $P(\mathbf{x}|y = j) = 0$  for all values of  $y$
  - How do we make predictions for such examples?
- To avoid overfitting
  - Make some bold assumptions to simplify the joint distribution

# The Naïve Bayes Assumption

- Assume that each feature is independent of any other features given the class label

$$\begin{aligned} &P(x_1 = u_1, \dots, x_m = u_m \mid y = j) \\ &= P(x_1 = u_1 \mid j) \cdots P(x_m = u_m \mid j) \end{aligned}$$



# A note about independence

- Assume A and B are two random events.  
Then

“A and B are independent”

if and only if

$$P(A|B) = P(A)$$

# Independence Theorems

- Assume  $P(A|B) = P(A)$
- Then

$$P(A \wedge B) = P(A) P(B)$$

- Assume  $P(A|B) = P(A)$
- Then

$$P(B|A) = P(B)$$

# Examples of independent events

- Two separate coin tosses
- Consider the following four events:
  - T: Toothache ( I have a toothache)
  - C: Catch (dentist's steel probe catches in my tooth)
  - A: Cavity (I have a cavity)
  - W: Weather (weather is good)
  - $P(T, C, A, W) = P(T, C, A)P(W)$

# Conditional Independence

- $P(A|B,C) = P(A|C)$ 
  - A and B are conditionally independent given C
  - $P(B|A,C) = P(B|C)$
- If A and B are conditionally independent given C, then we have
  - $P(A,B|C) = P(A|C) P(B|C)$

# Example of conditional independence

- T: Toothache ( I have a toothache)
- C: Catch (dentist's steel probe catches in my tooth)
- A: Cavity

T and C are conditionally independent given A:  $P(T, C|A) = P(T|A) * P(C|A)$

So , **events that are not independent from each other might be conditionally independent given some fact**

It can also happen the other way around. **Events that are independent might become conditionally dependent given some fact.**

B=Burglar in your house; A = Alarm (Burglar) rang in your house

E = Earthquake happened

B is independent of E (ignoring some minor possible connections between them)

However, if we know A is true, then B and E are no longer independent. Why?

$P(B|A) \gg P(B|A, E)$  Knowing E is true makes it much less likely for B to be true

# Naïve Bayes Classifier

- By assuming that each attribute is independent of any other attributes given the class label, we now have a *Naïve* Bayes Classifier
- Instead of learning a joint distribution of all features, we learn  $p(x_i|y = j)$  separately for each feature  $x_i$
- And we compute the joint by taking the product:

$$P(\mathbf{x} = [u_1, \dots, u_m]^T | y = j) = \prod_{i=1}^m P(x_i = u_i | y = j)$$

# Naïve Bayes Classifiers in a nutshell

1. Learn the  $p(x_i | y = j)$  for each feature  $x_i$ , and  $y$  value  $j$
  2. Estimate  $P(y = j)$  as fraction of records with  $y = j$ .
  3. For a new example  $x = [u_1, \dots, u_m]^T$ :
- } *learning*

$$\begin{aligned} y^{\text{predict}} &= \underset{j}{\operatorname{argmax}} P(y = j \mid x_1 = u_1, \dots, x_m = u_m) \\ &= \underset{j}{\operatorname{argmax}} P(x_1 = u_1, \dots, x_m = u_m \mid y = j) P(y = j) \\ &= \underset{j}{\operatorname{argmax}} P(y = j) \prod_{i=1}^m P(x_i = u_i \mid y = j) \end{aligned}$$

# Example

$X_1$	$X_2$	$X_3$	Y
1	1	1	0
1	1	0	0
0	0	0	0
0	1	0	1
0	0	1	1
0	1	1	1

Apply Naïve Bayes, and make prediction for (1,0,1)?



# When Bag-of-words uses counts (Multinomial Naïve Bayes)

- Often text is represented by the number of times each word appeared in it
  - E.g., “There are some shinning apples on the apple tree”
  - Word “apple” will be counted as appearing twice (ignoring the difference between plural and singular form)
  - How do we use naïve bayes in this case?
- We still learn  $p(w_i|y = 1)$  for each word  $i$
- It will be estimated as

$$\frac{\text{\# of times word } i \text{ appeared in spam emails}}{\text{total \# words in spam emails}}$$

- For the sentence S above, we will have:

$$\begin{aligned} P(S|y = 1) \\ = P(\text{"there"}|y = 1)P(\text{"are"}|y = 1) \dots P(\text{"apple"}|y = 1)^2 \dots P(\text{"tree"}|y = 1) \end{aligned}$$

Similarly for  $P(S|y = 0)$

- Bayes rule can then be used to compute  $P(y = 1|S)$

# Laplace Smoothing

- With Naïve Bayes Assumption, we still have zero probabilities
- E.g., if we receive an email that contains a word  $w$  that has never appeared in the training emails
  - $P(w|spam) = 0$  and  $P(w|nonspam) = 0$
- As such we ignore all the other words in the email because of this single rare word
- Laplace smoothing can help

Binary:

$$P(w|spam) = \frac{\text{\# of spam emails with word } w + 1}{\text{\# of spam emails} + 2}$$

Multinomial:

$$P(w|spam) = \frac{\text{\# of times word } w \text{ appeared in spam emails} + 1}{\text{total \# words in spam emails} + m}$$

# Final Notes about (Naïve) Bayes Classifier

- Any density estimator can be plugged in to estimate  $P(x_1, x_2, \dots, x_m | y)$  for Bayes, or  $P(x_i | y)$  for Naïve Bayes
- Real valued attributes can be discretized or directly modeled using simple continuous distributions such as Gaussian (Normal) distribution
- Naïve Bayes is wonderfully cheap and survives tens of thousands of attributes easily
- Laplace smoothing is important to avoid extreme probabilities