

Logistic regression

cs434

Logistic regression

- Recall the problem of regression
 - Learns a linear mapping from input vector \mathbf{x} to a continuous output $y \in (-\infty, \infty)$
- Logistic regression can be viewed as extending linear regression to handle binary output y by warping the output of a linear function to the range between 0 and 1
- For convenience, we will assume $y \in \{0,1\}$ for this lecture

Logistic regression (cont.)


- Consider the linear regression function

$$\mathbf{w}^T \mathbf{x} = w_0 + w_1 x_1 + \cdots + w_m x_m$$

- We introduce a function g :

$$g(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

Referred to as the sigmoid function or logistic function



Logistic Regression Makes Probabilistic Prediction

- We interpret the output of logistic regression probabilistically:

$$g(\mathbf{w}^T \mathbf{x}) = P(y = 1 | \mathbf{x}; \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

i.e., probability of $y = 1$ given the input \mathbf{x} and the model's parameter is \mathbf{w}

$$\begin{aligned} P(y = 0 | \mathbf{x}; \mathbf{w}) &= 1 - g(\mathbf{w}^T \mathbf{x}) \\ &= \frac{\exp(-\mathbf{w}^T \mathbf{x})}{1 + \exp(-\mathbf{w}^T \mathbf{x})} \end{aligned}$$

Logistic Regression forms a **linear** decision boundary

- We predict $y = 1$ if

$$p(y = 1|\mathbf{x}; \mathbf{w}) > p(y = 0|\mathbf{x}; \mathbf{w})$$

Predict $y = 1$ if

$$\frac{p(y = 1|\mathbf{x})}{p(y = 0|\mathbf{x})} > 1$$

Predict $y = 1$ if

$$\log \frac{p(y = 1|\mathbf{x})}{p(y = 0|\mathbf{x})} = w_0 + w_1x_1 + \dots + w_mx_m > 0$$

Odds of $y=1$

Side Note:

the odds of an event are the quantity $p / (1 - p)$, where p is the probability of the event

If I toss a fair dice, what are the odds that I will have a six?

LR assumes that the log odds of $y=1$ is a linear function of the input features

Learning \mathbf{w} for logistic regression

- Given a set of training data points (\mathbf{x}^i, y^i) , $i = 1, \dots, n$, the goal of learning is to find a weight vector \mathbf{w} such that

$$g(\mathbf{w}^T \mathbf{x}) = P(y = 1 | \mathbf{x}; \mathbf{w}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$

- is large (approaching 1) for positive ($y = 1$) training examples
- is small (approaching 0) for negative ($y = 0$) training examples

Learning Objective for Logistic Regression

- Consider a general form of objective (to be minimized):


$$L(\mathbf{w}) = \sum_{i=1}^n l(g(\mathbf{w}^T \mathbf{x}^i), y^i)$$

- For logistic regression the loss function is:

$$\begin{aligned} l(g(\mathbf{w}^T \mathbf{x}^i), y^i) \\ = \begin{cases} -\log g(\mathbf{w}^T \mathbf{x}^i) & \text{if } y^i = 1 \\ -\log(1 - g(\mathbf{w}^T \mathbf{x}^i)) & \text{if } y^i = 0 \end{cases} \end{aligned}$$

Loss function for logistic regression

$$l(g(\mathbf{w}^T \mathbf{x}^i), y^i) = \begin{cases} -\log P(y = 1 | x^i; w) & \text{if } y^i = 1 \\ -\log(1 - P(y = 1 | x^i; w)) & \text{if } y^i = 0 \end{cases}$$




When $y^i = 1$,

- if we predict $P(y = 1 | x^i; w) = 1$, the loss is 0
- If we predict $P(y = 1 | x^i; w) = 0$, the loss is ∞

Loss function for logistic regression

$$l(g_{\mathbf{w}}(x^i), y^i) = \begin{cases} -\log P(y = 1|x^i; \mathbf{w}) & \text{if } y^i = 1 \\ -\log(1 - P(y = 1|x^i; \mathbf{w})) & \text{if } y^i = 0 \end{cases}$$



When $y^i = 0$,

- if we predict $P(y = 1|x^i; \mathbf{w}) = 0$, the loss is 0
- If we predict $P(y = 1|x^i; \mathbf{w}) = 1$, the loss is ∞

Representing it compactly

$$l(g_{\mathbf{w}}(x^i), y^i) = \begin{cases} -\log P(y = 1|x^i; w) & \text{if } y^i = 1 \\ -\log(1 - P(y = 1|x^i; w)) & \text{if } y^i = 0 \end{cases}$$


can be represented compactly as:

$$l(g_{\mathbf{w}}(x^i), y^i) = -y^i \log P(y = 1|x^i; w) - (1 - y^i) \log(1 - P(y = 1|x^i; w))$$

Minimizing $L(\mathbf{w})$

- Unfortunately this does not have a close form solution
 - You take the derivative, set it to zero, but no closed form solution
- Instead, we iteratively search for the optimal \mathbf{w}
- Start with a random \mathbf{w} , iteratively improve \mathbf{w} (similar to Perceptron) by taking the negative gradient
- This is referred to as gradient descent

Gradient descent to minimize $L(\mathbf{w})$



1. Start from some initial guess \mathbf{w}^0
2. Find the direction of steepest descent – opposite of the gradient direction $\nabla f(\mathbf{w})$
3. Take a step toward that direction
$$\mathbf{w}^{t+1} = \mathbf{w}^t - \nabla f(\mathbf{w}^t)$$
4. Repeat until no local improvement is possible
($|\nabla f(\mathbf{w}^t)| \leq \epsilon$)

Different starting points may lead to different local minimum if objective is not convex (like what's shown in the above figure)

Gradient of $L(\mathbf{w})$

$$l(g(\mathbf{w}^T \mathbf{x}^i), y^i) = -y^i \log P(y = 1 | x^i; w) - (1 - y^i) \log(1 - P(y = 1 | x^i; w))$$

$$l(g(\mathbf{w}^T \mathbf{x}^i), y^i)$$

$$= -y^i \log g(\mathbf{w}^T \mathbf{x}^i) - (1 - y^i) \log(1 - g(\mathbf{w}^T \mathbf{x}^i))$$

Useful fact: $g'(t) = t(1 - t)$

$$\nabla g(\mathbf{w}^T \mathbf{x}^i) = g(\mathbf{w}^T \mathbf{x}^i) (1 - g(\mathbf{w}^T \mathbf{x}^i)) \mathbf{x}^i$$

$$\nabla l(g(\mathbf{w}^T \mathbf{x}^i), y^i) = (y^i - g(\mathbf{w}^T \mathbf{x}^i)) \mathbf{x}^i$$

Online gradient descent for Logistic Regression

Note: y takes 0/1 here, not 1/-1

Given : training examples (\mathbf{x}^i, y^i) , $i = 1, \dots, N$

Let $\mathbf{w} \leftarrow (0, 0, 0, \dots, 0)$

Repeat until convergence

For every example i

$$\hat{y}^i \leftarrow \frac{1}{1 + e^{-\mathbf{w}^T \mathbf{x}^i}}$$

$$\mathbf{w} \leftarrow \mathbf{w} + \eta \cdot (y^i - \hat{y}^i) \cdot \mathbf{x}^i$$

η is the learning rate or step size

Note the striking similarity between LR and perceptron
Both learn a linear decision boundary
The iterative algorithm takes very similar form

Daja vu?

Batch Learning for Logistic Regression

Note: y takes 0/1 here, not 1/-1

Given : training examples (\mathbf{x}^i, y^i) , $i = 1, \dots, N$

Let $\mathbf{w} \leftarrow (0, 0, 0, \dots, 0)$

Repeat until convergence

$\mathbf{d} \leftarrow (0, 0, 0, \dots, 0)$

For $i = 1$ to N do

$$\hat{y}^i \leftarrow \frac{1}{1 + e^{-\mathbf{w} \cdot \mathbf{x}^i}}$$

$$error = y^i - \hat{y}^i$$

$$\mathbf{d} = \mathbf{d} + error \cdot \mathbf{x}^i$$

$$\mathbf{w} \leftarrow \mathbf{w} + \eta \mathbf{d}$$

Logistic regression: Summary

- LR uses the logistic function to warp the output of a linear function to between zero and one, which is interpreted as $P(y = 1|\mathbf{x}; \mathbf{w})$
- Learning aims to learn a vector \mathbf{w} s.t. examples with $y = 1$ are predicted to have high $P(y = 1|\mathbf{x}; \mathbf{w})$ and vice versa
- Learn \mathbf{w} iteratively using gradient descent
- Strong similarity with Perceptron
- Logistic regression learns a linear decision boundaries
 - By introducing nonlinear features (i.e., $x_1^2, x_2^2, x_1x_2, \dots$), can be extended to nonlinear boundary.